

1. move "every claim has a source" up to the hero. it's the strongest thing you have and right now it's buried in a bullet way down the page. it's what actually separates you from someone just prompting chatgpt.
2. kill the "vs. prompting an AI" framing. the section header ("what makes it different from just prompting an ai") concedes the fight before it starts. rename to something like "what every castor report guarantees" or "four things you can count on."
3. the viability score needs a methodology page, or drop the number. a 1-100 score with no explanation of how it's calculated will get eye-rolled by exactly the people you're targeting. founders and VCs are the two groups most allergic to fake precision.
4. put a screenshot of a real report above the fold. right now the sample is a click away. it should be visible.
5. pick a lane: fast or rigorous. "runs in minutes" and "research-grade" pull in opposite directions. for a \$29 buyer, lean into fast and directional. save "rigorous" for the enterprise page.
6. this line is a little smug: "you do not need to know the right questions to ask. the system does." just say "the chat asks you what it needs."
7. no proof anywhere on the page. no logos, testimonials, or "we've generated X reports." even a "used by teams at Y and Z" line lowers the risk of clicking buy.
8. the 5-pack pricing math is soft. \$29 to \$20 per report is only about 30% off. either widen the gap (bump single to \$39, keep 5-pack at \$99) or drop the pack entirely.